Steven Pan (span3@ucsc.edu)
CMPS 111: Intro to Operating Systems
Project 3: Due 2/21/14

Memory Management Report

Preface: Due to being unable to implement random fit, I had to take a guess for its stats.

Best fit: Mostly very small holes, a few extremely large holes left open in between.

Worst fit: Many large holes, very few small holes.

Random fit: Roughly equal amounts of small and large holes.

Experiment:

Since buddy allocation has to request another equally large section, it would work optimally if a process were to request memory at or extremely close to a power of 2. If the workload also involved chunks of memory close to a power of 2, buddy allocation would leave very little memory holes and would be extremely efficient.